

# Lecture 4: Internetworking Protocols

- » Application Layer  
(contd).

# User-server state: cookies

Many major Web sites use cookies

## Four components:

- 1) cookie header line of HTTP response message
- 2) cookie header line in HTTP request message
- 3) cookie file kept on user's host, managed by user's browser
- 4) back-end database at Web site


## Example:

- r Susan always access Internet always from PC
- r visits specific e-commerce site for first time
- r when initial HTTP requests arrives at site, site creates:
  - ❖ unique ID
  - ❖ entry in backend database for ID

# Cookies: keeping "state" (cont.)

client

server



cookie file



one week later:



usual http request msg

usual http response  
**Set-cookie: 1678**

usual http request msg  
**cookie: 1678**

usual http response msg

usual http request msg  
**cookie: 1678**

usual http response msg

Amazon server  
creates ID  
1678 for user


create  
entry

cookie-  
specific  
action

cookie-  
specific  
action

access

access



backend  
database

# Cookies (continued)

## What cookies can bring:

- ❑ authorization
- ❑ shopping carts
- ❑ recommendations
- ❑ user session state (Web e-mail)

## How to keep "state":

- r protocol endpoints: maintain state at sender/receiver over multiple transactions
- r cookies: http messages carry state

— aside —


### Cookies and privacy:

- r cookies permit sites to learn a lot about you
- r you may supply name and e-mail to sites

# Web caches (proxy server)

**Goal:** satisfy client request without involving origin server

- r user sets browser: Web accesses via cache
- r browser sends all HTTP requests to cache
  - ❖ object in cache: cache returns object
  - ❖ else cache requests object from origin server, then returns object to client



# More about Web caching

- r cache acts as both client and server

- r typically cache is installed by ISP (university, company, residential ISP)

## Why Web caching?

- r reduce response time for client request

- r reduce traffic on an institution's access link.

- r Internet dense with caches: enables "poor" content providers to effectively deliver content (but so does


# Caching example

## Assumptions

- r average object size = 1,000,000 bits
- r avg. request rate from institution's browsers to origin servers = 15/sec
- r delay from institutional router to any origin server and back to router = 2 sec

## Consequences

- r utilization on LAN = 15%
- r utilization on access link = 100%
- r total delay = Internet delay + access delay + LAN delay  
= 2 sec + minutes + milliseconds




# Caching example (cont)

## possible solution

- r increase bandwidth of access link to, say, 100 Mbps

## consequence

- r utilization on LAN = 15%
- r utilization on access link = 15%
- r Total delay = Internet delay + access delay + LAN delay  
= 2 sec + msecs + msecs
- r often a costly upgrade






# Caching example (cont)

## possible solution: install cache

r suppose hit rate is 0.4

## consequence

- r 40% requests will be satisfied almost immediately
- r 60% requests satisfied by origin server
- r utilization of access link reduced to 60%, resulting in negligible delays (say 10 msec)
- r total avg delay = Internet delay + access delay + LAN delay  
$$= .6 \cdot (2.01) \text{ secs} + .4 \cdot \text{milliseconds} < 1.4 \text{ secs}$$



# Conditional GET

**Goal:** don't send object if cache

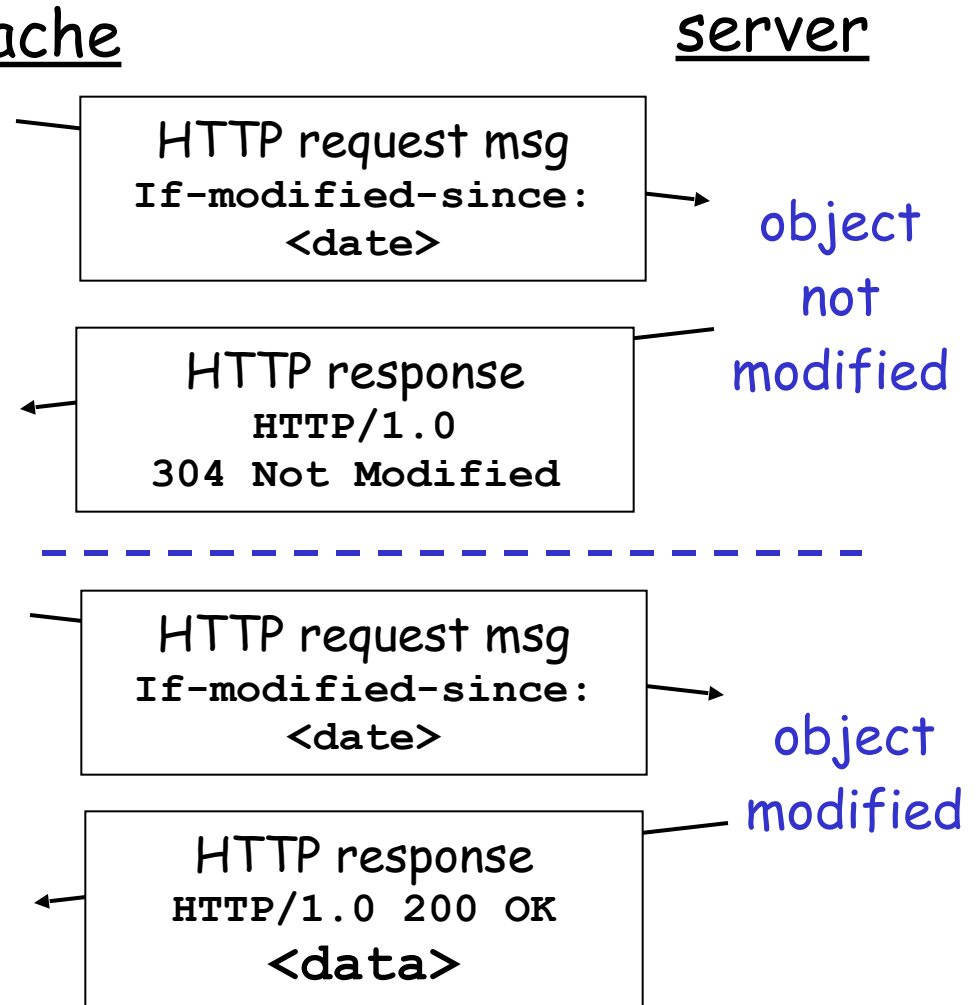
has up-to-date cached version

cache: specify date of cached copy in HTTP request

If-modified-since: <date>

server: response contains no object if cached copy is up-to-date:


HTTP/1.0 304 Not Modified



# Chapter 2: Application layer

- r 2.1 Principles of network applications
- r 2.2 Web and HTTP
- r 2.3 FTP
- r 2.4 Electronic Mail
  - ❖ SMTP, POP3, IMAP
- r 2.5 DNS
- r 2.6 P2P applications
- r 2.7 Socket programming with UDP
- r 2.8 Socket programming with TCP


# FTP: the file transfer protocol



- r transfer file to/from remote host
- r client/server model
  - ❖ **client**: side that initiates transfer (either to/from remote)
  - ❖ **server**: remote host
- r ftp: RFC 959
- r ftp server: port 21

# FTP: separate control, data connections

- r FTP client contacts FTP server at port 21, TCP is transport protocol
- r client authorized over control connection
- r client browses remote directory by sending commands over control connection.
- r when server receives file transfer command, server opens 2<sup>nd</sup> TCP connection (for file) to client
- r after transferring one file, server closes data connection.



- r server opens another TCP data connection to transfer another file.
- r control connection: "out of band"
- r FTP server maintains "state": current directory, earlier authentication

# FTP commands, responses

## Sample commands:

- r sent as ASCII text over control channel
- r `USER username`
- r `PASS password`
- r `LIST` return list of file in current directory
- r `RETR filename` retrieves (gets) file
- r `STOR filename` stores (puts) file onto remote host

## Sample return codes

- r status code and phrase (as in HTTP)
- r 331 Username OK, password required
- r 125 data connection already open; transfer starting
- r 425 Can't open data connection
- r 452 Error writing file

# Chapter 2: Application layer

- r 2.1 Principles of network applications
- r 2.2 Web and HTTP
- r 2.3 FTP
- r 2.4 Electronic Mail
  - ❖ SMTP, POP3, IMAP
- r 2.5 DNS
- r 2.6 P2P applications
- r 2.7 Socket programming with UDP
- r 2.8 Socket programming with TCP

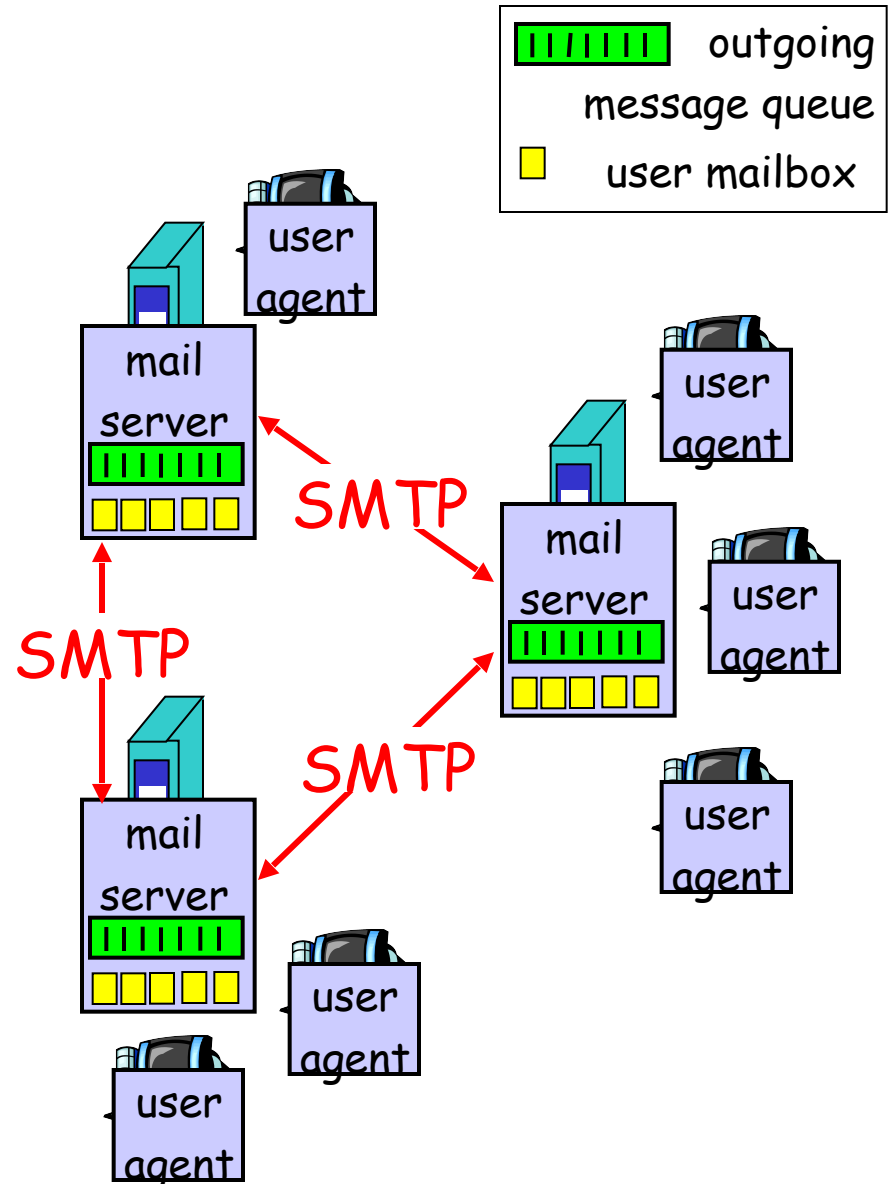
# Electronic Mail

## Three major components:

- r user agents
- r mail servers
- r simple mail transfer protocol: SMTP

## User Agent

- r a.k.a. "mail reader"
- r composing, editing, reading mail messages
- r e.g., Eudora, Outlook, elm, Mozilla Thunderbird
- r outgoing, incoming messages stored on server

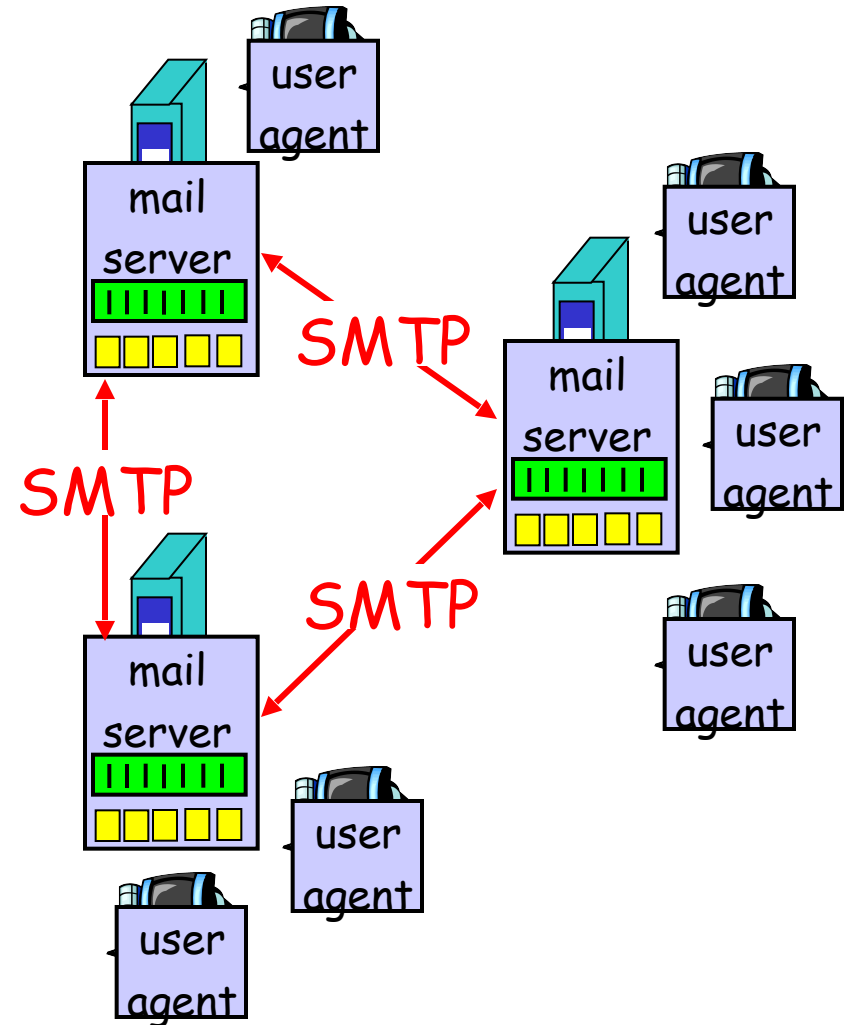




# Electronic Mail: mail servers

## Mail Servers

- r **mailbox** contains incoming messages for user
- r **message queue** of outgoing (to be sent) mail messages
- r **SMTP protocol** between mail servers to send email messages
  - ❖ client: sending mail server
  - ❖ "server": receiving mail server




# Electronic Mail: SMTP [RFC 2821]

- r uses TCP to reliably transfer email message from client to server, port 25
- r direct transfer: sending server to receiving server
- r three phases of transfer
  - ❖ handshaking (greeting)
  - ❖ transfer of messages
  - ❖ closure
- r command/response interaction
  - ❖ **commands**: ASCII text
  - ❖ **response**: status code and phrase
- r messages must be in 7-bit ASCII

# Scenario: Alice sends message to Bob

- 1) Alice uses UA to compose message and "to"  
`bob@someschool.edu`
- 2) Alice's UA sends message to her mail server; message placed in message queue
- 3) Client side of SMTP opens TCP connection with Bob's mail server
- 4) SMTP client sends Alice's message over the TCP connection
- 5) Bob's mail server places the message in Bob's mailbox
- 6) Bob invokes his user agent to read message



# Sample SMTP interaction

```
S: 220 hamburger.edu
C: HELO crepes.fr
S: 250 Hello crepes.fr, pleased to meet you
C: MAIL FROM: <alice@crepes.fr>
S: 250 alice@crepes.fr... Sender ok
C: RCPT TO: <bob@hamburger.edu>
S: 250 bob@hamburger.edu ... Recipient ok
C: DATA
S: 354 Enter mail, end with "." on a line by itself
C: Do you like ketchup?
C: How about pickles?
C: .
S: 250 Message accepted for delivery
C: QUIT
S: 221 hamburger.edu closing connection
```

## Try SMTP interaction for yourself:

- r `telnet servername 25`
- r see 220 reply from server
- r enter HELO, MAIL FROM, RCPT TO, DATA, QUIT commands

above lets you send email without using email client (reader)

# SMTP: final words

- r SMTP uses persistent connections
- r SMTP requires message (header & body) to be in 7-bit ASCII
- r SMTP server uses CRLF.CRLF to determine end of message

## Comparison with HTTP:

- r HTTP: pull
- r SMTP: push
- r both have ASCII command/response interaction, status codes
- r HTTP: each object encapsulated in its own response msg
- r SMTP: multiple objects sent in multipart msg

# Mail message format

SMTP: protocol for exchanging email msgs

RFC 822: standard for text message format:


r header lines, e.g.,

- ❖ To:
- ❖ From:
- ❖ Subject:

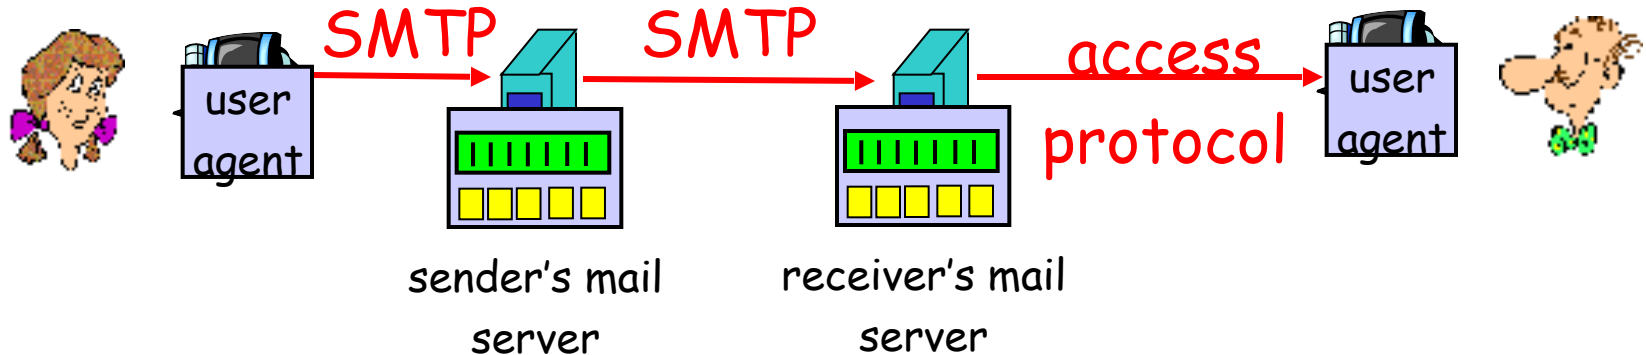
different from SMTP commands!

r body

- ❖ the "message", ASCII characters only



# Mail access protocols



- r SMTP: delivery/storage to receiver's server
- r Mail access protocol: retrieval from server
  - ❖ POP: Post Office Protocol [RFC 1939]
    - authorization (agent <-->server) and download
  - ❖ IMAP: Internet Mail Access Protocol [RFC 1730]
    - more features (more complex)
    - manipulation of stored msgs on server
  - ❖ HTTP: gmail, Hotmail, Yahoo! Mail, etc.



# POP3 protocol

## authorization phase

r client commands:

- ❖ user: declare username
- ❖ pass: password

r server responses

- ❖ +OK
- ❖ -ERR

## transaction phase, client:

r list: list message numbers

r retr: retrieve message by number

r dele: delete

r quit

```
S: +OK POP3 server ready
C: user bob
S: +OK
C: pass hungry
S: +OK user successfully logged on
```

```
C: list
S: 1 498
S: 2 912
S: .
C: retr 1
S: <message 1 contents>
S: .
C: dele 1
C: retr 2
S: <message 1 contents>
S: .
C: dele 2
C: quit
S: +OK POP3 server signing off
```

# POP3 (more) and IMAP

## More about POP3

- r Previous example uses "download and delete" mode.
- r Bob cannot re-read e-mail if he changes client
- r "Download-and-keep": copies of messages on different clients
- r POP3 is stateless across sessions

## IMAP

- r Keep all messages in one place: the server
- r Allows user to organize messages in folders
- r IMAP keeps user state across sessions:
- ❖ names of folders and mappings between message IDs and folder name

# Chapter 2: Application layer

- r 2.1 Principles of network applications
- r 2.2 Web and HTTP
- r 2.3 FTP
- r 2.4 Electronic Mail
  - ❖ SMTP, POP3, IMAP
- r 2.5 DNS
- r 2.6 P2P applications
- r 2.7 Socket programming with UDP
- r 2.8 Socket programming with TCP

# DNS (Domain Name Service)

## DNS services


- r hostname to IP address translation
- r host aliasing
  - ❖ Canonical, alias names
- r mail server aliasing
- r load distribution
  - ❖ replicated Web servers: set of IP addresses for one canonical name

## Why not centralize DNS?

- r single point of failure
- r traffic volume
- r distant centralized database
- r maintenance

doesn't scale!

# Distributed, Hierarchical Database



Client wants IP for www.amazon.com; 1<sup>st</sup> approx:

- r client queries a root server to find com DNS server
- r client queries com DNS server to get amazon.com DNS server
- r client queries amazon.com DNS server to get IP address for www.amazon.com